

ISTANBUL AYDIN UNIVERSITY
THE FACULTY OF ENGINEERING



HALIT AYDIN CAMPUS

FALL – 2020

COM447

OPEN SOURCE SOFTWARES

SIMPLE FILE SYSTEM

Ant Doğa ŞENTÜRK - B1605.010010

Selim CESUR - B1605.010003

DEPARTMENT : Computer Engineering

LECTURER: Dr. Öğr. Üyesi ADEM ÖZYAVAŞ

Introduction

In this project, we created a simple file system. File system is composed of three main parts, super block, inode table and data block table. We create this file system with the name sfs.bin. This file contains the metadata and the actual data for the file system. We can use five different commands in the file system. “mkdir” command to create directory, “mkfile” command to create a regular file, “cd” command to change directory, “ls” command to print the list of files in the current directory and finally “lsrec” command to list the files in the whole file system.

Code:

We have 5 different global variables in this program. Two of them are to decide the type of the file we create, one is for the size of a datablock, one is for our Simple File System file and the last one is to hold the inode number of the current directory.

```
6  const int REG_FILE = 0;
7  const int DIRECTORY = 1;
8  const int DATA_BLOCK_SIZE = 512;
9
10 FILE *sfs;
11 int currentDirectory = 0;
```

We have three structures for the three main parts of our file system. First one is the super block. Super block holds the information of which indices of the inode table and which data blocks are used by our file system. We want to be able to have at most 32 files in our file system, for that reason, we use integer as an inode bitmap. Furthermore, we want to be able to allocate at most 10 data blocks for every file in the system. That's why we use integer array with 10 elements for data bitmap.

Second structure is for the directories or regular files in the file system. It holds the type information, size information and the numbers of the data blocks that file uses.

The last structure is for directory entries. This is where we give name to the files and see in which file they are in by looking at their inode number.

```
22 struct super_block {
23     int inode_bitmap;
24     int data_bitmap[10];
25 };
26
27 struct inode_st {
28     int type;
29     int size;
30     int data_block_indices[10];
31 };
32
33 struct dir_ent {
34     char name [28];
35     unsigned int inode_no;
36 };
```

We also create the super block as a global variable since we will use it in many functions.

```
42 struct super_block sb;
```

When we run the program, first the sfs file is gets created. Then, we initializ the file system, then we run the simpleBash() function which is an infinite loop that reads the commands.

```
44 int main() {
45     sfs = fopen("C:\\Users\\selim\\Desktop\\File System\\sfs.bin", "w+");
46     init_fileSystem();
47
48     simpleBash();
49
50     return 0;
51 }
```

When we initialize the file system, we first make everything in the super block 0, because there is no files in the file system yet.

Then we create the first inode named “root”. Everything in the file system will be inside this root. We set the type of the root directory and we set it’s size 0, since there is nothing inside it yet. Lastly, we allocate the first data block for root.

Now that the root uses the first inode in the inode table and first data block in the data block table, we make the first bit of the relevant bitmaps 1.

```
309 void init_fileSystem() {
310     sb.inode_bitmap = 0;
311
312     int i;
313     for(i = 0; i < 10; i++)
314         sb.data_bitmap[i] = 0;
315
316     struct inode_st root;
317     root.type = DIRECTORY;
318     root.size = 0;
319
320     for(i = 0; i < 10; i++)
321         root.data_block_indices[i] = 0;
322
323     sb.inode_bitmap = 1;
324     sb.data_bitmap[0] = 1;
```

Every directory will have two directories inside it (. and ..) as default. We create those two directories an entries. Since they are inside the root and the inode number for the root is 0, we set their inode number values to 0. After creating two entries, we add their size to the root’s size.

Now that we set the values of everything related to super block and the root, we write it to the sfs file. We use fwrite command to write the super block, this operation writes the super block to the beginning of the file. Then we write the root next to it using same command.

In order to write the two entries . and .. we jump to the beginning of the data blocks with fseek command. Since root is using the first data block, we write the entries to the beginning of the data block with fwrite command.

```

326     struct dir_ent dot;
327     strcpy(dot.name, ".");
328     dot.inode_no = 0;
329
330     struct dir_ent dotdot;
331     strcpy(dotdot.name, "..");
332     dotdot.inode_no = 0;
333
334     root.size = sizeof(struct dir_ent)*2;
335
336     fwrite(&sb, sizeof(struct super_block), 1, sfs);
337     fwrite(&root, sizeof(struct inode_st), 1, sfs);
338
339     fseek(sfs, sizeof(struct super_block)+(32*sizeof(struct inode_st))+(root.data_block_indices[9]*DATA_BLOCK_SIZE), SEEK_SET);
340
341     fwrite(&dot, sizeof(struct dir_ent), 1, sfs);
342     fwrite(&dotdot, sizeof(struct dir_ent), 1, sfs);
343 }

```

After initializing the file system. We run our bash, which is an infinite loop. Inside the loop, we take the command as an input from the user. First command is “mkdir”. If the input from the user is mkdir <directory name>, we compare this directory name with the names of all the files in the current directory. If a file with the same name doesn’t exist, we run the mkdir() function to create a new directory.

```

54 void simpleBash() {
55     printf("Write 'help' to see the possible commands.\n");
56     char command[32];
57     while(1) {
58         printf("> ");
59         scanf("%s", command);
60         if(strcmp(command, "mkdir") == 0) {
61             bool fileExists = false;
62             char fileName[32];
63             scanf("%s", fileName);
64             struct inode_st tempInode = getCurrentInode();
65
66             int numOfEntries = tempInode.size/sizeof(struct dir_ent);
67             fseek(sfs, sizeof(struct super_block)+32*sizeof(struct inode_st)+(tempInode.data_block_indices[9]*DATA_BLOCK_SIZE), SEEK_SET);
68
69             struct dir_ent tempEntry;
70             int i;
71             for(i=0; i<numOfEntries; i++){
72                 fread(&tempEntry, sizeof(struct dir_ent), 1, sfs);
73                 if(strcmp(tempEntry.name, fileName) == 0) {
74                     printf("File with the same name already exists!\n");
75                     fileExists = true;
76                     break;
77                 }
78             }
79             if(!fileExists) {
80                 mkdir(fileName);
81             }
82         }
83     }
84 }

```

In the mkdir() function, we first create a new inode structure called newDir. We set the size of it to 0 initially. Then, we use getInodeNo() and getDataBlockNo() functions to return the first empty slot in the inode table and in the data block table.

```

130 void mkdir(char *fileName) {
131     struct inode_st newDir;
132     newDir.type = DIRECTORY;
133     newDir.size = 0;
134
135     int inode = getInodeNo(sb);
136     int datablock = getDataBlockNo(sb);

```

How we get the first free slot in the inode table is, we do an operation between the inode_bitmap and 1. This operation returns 0 if the least significant bit of the inode_bitmap is 0. If not, we right shift the inode_bitmap by 1 to change the least significant bit of it. Then compare it again until we get the 0. Getting the data block number works the same way.

```

356 int getInodeNo(struct super_block sb) {
357     int i;
358     for(i=0; i<32; i++) {
359         if((sb.inode_bitmap & 1) == 0) {
360             return i;
361         }
362         else {
363             sb.inode_bitmap = sb.inode_bitmap >> 1;
364         }
365     }
366     printf("There is not enough space in the file system!");
367     exit(0);
368 }
369
370 int getDataBlockNo(struct super_block sb) {
371     int i;
372     int j;
373     for(i=0; i<10; i++) {
374         for(j=0; j<32; j++) {
375             if((sb.data_bitmap[i] & 1) == 0) {
376                 return (i*32)+j;
377             }
378             else {
379                 sb.data_bitmap[i] = sb.data_bitmap[i] >> 1;
380             }
381         }
382     }
383     printf("There is not enough space in the file system!");
384     exit(0);
385 }

```

After we get the inode number and data block number, we set the data block indices for the new directory, and create the (.) and (..) entries same with the root. Since we add the directory entries to new directory, we increase the size of the accordingly.

```

138     int i;
139     for(i=0; i<10; i++)
140         newDir.data_block_indices[i] = datablock;
141
142     struct dir_ent dot;
143     strcpy(dot.name, ".");
144     dot.inode_no = inode;
145
146     struct dir_ent dotdot;
147     strcpy(dotdot.name, "..");
148     dotdot.inode_no = currentDirectory;
149
150     newDir.size = sizeof(struct dir_ent)*2;

```

Now that we use the inode and data block that we get for our new directory, we set them to 1 in the relative bitmaps. After updating the bitmaps, we update the super block by writing it to file again.

Another thing we need to update is the size of the current directory because we will add the new directory as a directory entry in the current directory. We add it's size another size of a director entry and then re-write it to the file.

```

152     int mask = 1 << inode;
153     sb.inode_bitmap = (sb.inode_bitmap & ~mask) | ((1 << inode) & mask);
154
155     int index = datablock/32;
156     int p = datablock%32;
157     mask = 1 << p;
158     sb.data_bitmap[index] = (sb.data_bitmap[index] & ~mask) | ((1 << p) & mask);
159
160     fseek(sfs, 0, SEEK_SET);
161     fwrite(&sb, sizeof(struct super_block), 1, sfs);
162
163     struct inode_st tempInode = getCurrentInode();
164
165     tempInode.size += sizeof(struct dir_ent);
166
167     fseek(sfs, sizeof(struct super_block)+(currentDirectory*sizeof(struct inode_st)), SEEK_SET);
168     fwrite(&tempInode, sizeof(struct inode_st), 1, sfs);

```

Now, we jump to the correct place in the inode table and write the new directory there. Then, we jump to the data block table and write the two entries (.) and (..).

Lastly, we create a directory entry for the new directory and add it to the data block of the current directory.

```

170     fseek(sfs, sizeof(struct super_block)+inode*sizeof(struct inode_st), SEEK_SET);
171     fwrite(&newDir, sizeof(struct inode_st), 1, sfs);
172
173     fseek(sfs, sizeof(struct super_block)+(32*sizeof(struct inode_st))+(newDir.data_block_indices[9]*DATA_BLOCK_SIZE), SEEK_SET);
174     fwrite(&dot, sizeof(struct dir_ent), 1, sfs);
175     fwrite(&dotdot, sizeof(struct dir_ent), 1, sfs);
176
177     struct dir_ent newDirEnt;
178     strcpy(newDirEnt.name, fileName);
179     newDirEnt.inode_no = inode;
180
181     fseek(sfs, sizeof(struct super_block)+(32*sizeof(struct inode_st))+(tempInode.data_block_indices[9]*DATA_BLOCK_SIZE)+
182         (tempInode.size-32), SEEK_SET);
183
184     fwrite(&newDirEnt, sizeof(struct dir_ent), 1, sfs);
185 }

```

Second command is “mkfile”. This command works exactly the same way with mkdir command. Only difference is that we don’t add (.) and (..) entries to it’s data block. We instead add a text to it. Then, we set the size of it to the size of the text.

```
186 void mkfile(char *fileName) {
187
188     char text[] = "My name is Selim and his name is Doğa. We have met at English preparation school 4 years ago...";
189
190     struct inode_st newFile;
191     newFile.type = REG_FILE;
192     newFile.size = sizeof(text);
```

Third command is “cd”. We use this command to change the current Directory. First, we get the current Directory and see how many files inside it by dividing it’s size by the size of an entry.

```
233 void cd() {
234     char fileName[32];
235     scanf("%s", fileName);
236     bool fileExists = false;
237
238     struct inode_st tempInode = getCurrentInode();
239
240     int numOfEntries = tempInode.size/sizeof(struct dir_ent);
241     fseek(sfs, sizeof(struct super_block)+32*sizeof(struct inode_st)+(tempInode.data_block_indices[9]*DATA_BLOCK_SIZE), SEEK_SET);
```

Then, we iterate through all files and see the directory we want to go does exist. If we find the file, we check it’s type. If it is a directory, we change the global variable current Directory to it’s inode number. If it’s not a directory, we just send a informative message and don’t do anything else. If we don’t find the file, we again print a message to inform the user that the file doesn’t exist.

```
243     struct dir_ent tempEntry;
244     int i;
245     for(i=0; i<numOfEntries; i++){
246         fread(&tempEntry, sizeof(struct dir_ent), 1, sfs);
247         if(strcmp(tempEntry.name, fileName) == 0) {
248
249             fileExists = true;
250             fseek(sfs, sizeof(struct super_block)+(tempEntry.inode_no*sizeof(struct inode_st)), SEEK_SET);
251
252             struct inode_st tempInode;
253             fread(&tempInode, sizeof(struct inode_st), 1, sfs);
254
255             if(tempInode.type == DIRECTORY) {
256                 currentDirectory = tempEntry.inode_no;
257             }
258             else {
259                 printf("This file is not a directory!\n");
260             }
261             break;
262         }
263     }
264     if(!fileExists) {
265         printf("No such directory!\n");
266     }
```


Forth command is “ls”. We use ls to print the list of the files in the current Directory. We get the current directory and iterate all the files inside of it’s data block the same way we did in the cd command. This time, while we iterate them, we print them to the console.

```

269 void ls() {
270     struct inode_st tempInode = getCurrentInode();
271
272     int numOfEntries = tempInode.size/sizeof(struct dir_ent);
273     fseek(sfs, sizeof(struct super_block)+32*sizeof(struct inode_st)+(tempInode.data_block_indices[9]*DATA_BLOCK_SIZE), SEEK_SET);
274
275     struct dir_ent tempEntry;
276     int i;
277     for(i=0; i<numOfEntries; i++){
278         fread(&tempEntry, sizeof(struct dir_ent), 1, sfs);
279         printf("Directory %d: %s\n", i+1, tempEntry.name);
280     }
281 }

```

The last command is “lsrec”. This command -differently from the “ls”- prints the whole file system when we write it. Same with the ls command, we get the current Directory and iterate the files in it’s data block. Instead of printing the file names, we store them in an array to print it later. After storing them, we iterate them again but this time we check their name and see if it’s a directory that we need go inside (directories that are not (.) and (..)). If they are, we change the local current directory variable and run the lsrec function with as this current directory variable is one of it’s arguments. This way lsrec becomes a recursive function that prints the whole file system.

```

283 void lsrec(int tab, int currentDirectory) {
284     struct inode_st tempInode = getCurrentInode();
285
286     if(tempInode.type == DIRECTORY) {
287         int numOfEntries = tempInode.size/sizeof(struct dir_ent);
288         fseek(sfs, sizeof(struct super_block)+32*sizeof(struct inode_st)+(tempInode.data_block_indices[9]*DATA_BLOCK_SIZE), SEEK_SET);
289
290         int i;
291         struct dir_ent entries[numOfEntries];
292         for(i=0; i<numOfEntries; i++){
293             fread(&entries[i], sizeof(struct dir_ent), 1, sfs);
294         }
295
296         int j;
297         for(i=0; i<numOfEntries; i++){
298             for(j=0; j<tab; j++) {
299                 printf("\t");
300                 printf(" ");
301             }
302             printf("Directory %d: %s\n", i+1, entries[i].name);
303             if(strcmp(entries[i].name, ".") != 0 && strcmp(entries[i].name, "..") != 0) {
304                 currentDirectory = entries[i].inode_no;
305                 tab++;
306                 lsrec(tab, currentDirectory);
307                 tab--;
308             }
309         }
310     }
311 }

```

This is the function that we used all the commands above to get current directory’s inode:

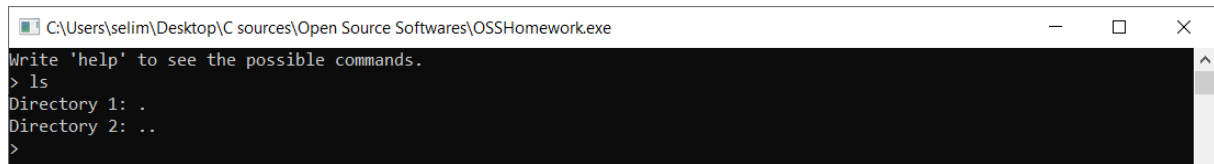
```

389 struct inode_st getCurrentInode() {
390     fseek(sfs, sizeof(struct super_block)+(currentDirectory*sizeof(struct inode_st)), SEEK_SET);
391     struct inode_st tempInode;
392     fread(&tempInode, sizeof(struct inode_st), 1, sfs);
393
394     return tempInode;
395 }

```

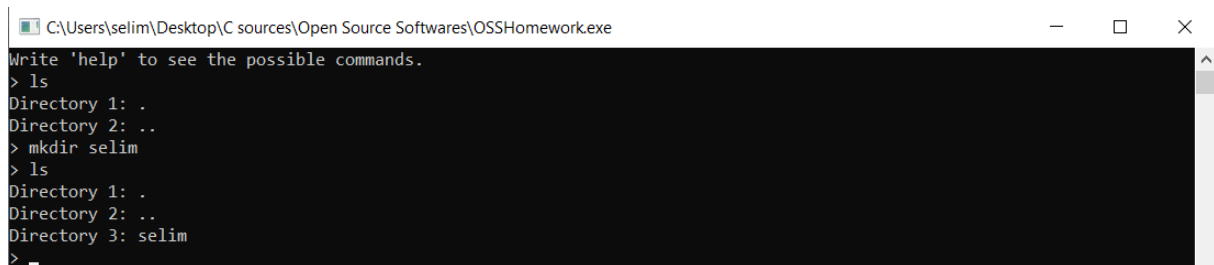

Sample Run:

First, we write “ls” to see that there only two directories (.) and (..) in the root.

A screenshot of a terminal window titled "C:\Users\selim\Desktop\C sources\Open Source Softwares\OSSHomework.exe". The terminal shows the prompt "Write 'help' to see the possible commands." followed by the command "> ls". The output is "Directory 1: ." and "Directory 2: ..". The prompt ">" is shown again.

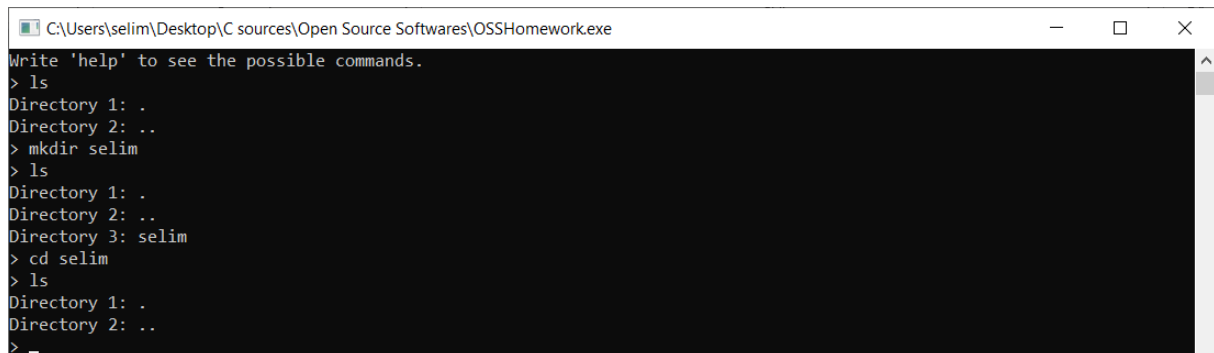
```
C:\Users\selim\Desktop\C sources\Open Source Softwares\OSSHomework.exe
Write 'help' to see the possible commands.
> ls
Directory 1: .
Directory 2: ..
>
```

Then, we run “mkdir” command to create a directory named “selim”. If we check with “ls” command, we see that it is created correctly.

A screenshot of a terminal window titled "C:\Users\selim\Desktop\C sources\Open Source Softwares\OSSHomework.exe". The terminal shows the prompt "Write 'help' to see the possible commands." followed by the command "> ls". The output is "Directory 1: ." and "Directory 2: ..". Then, the command "> mkdir selim" is entered, followed by "> ls". The output now includes "Directory 3: selim". The prompt ">" is shown again.

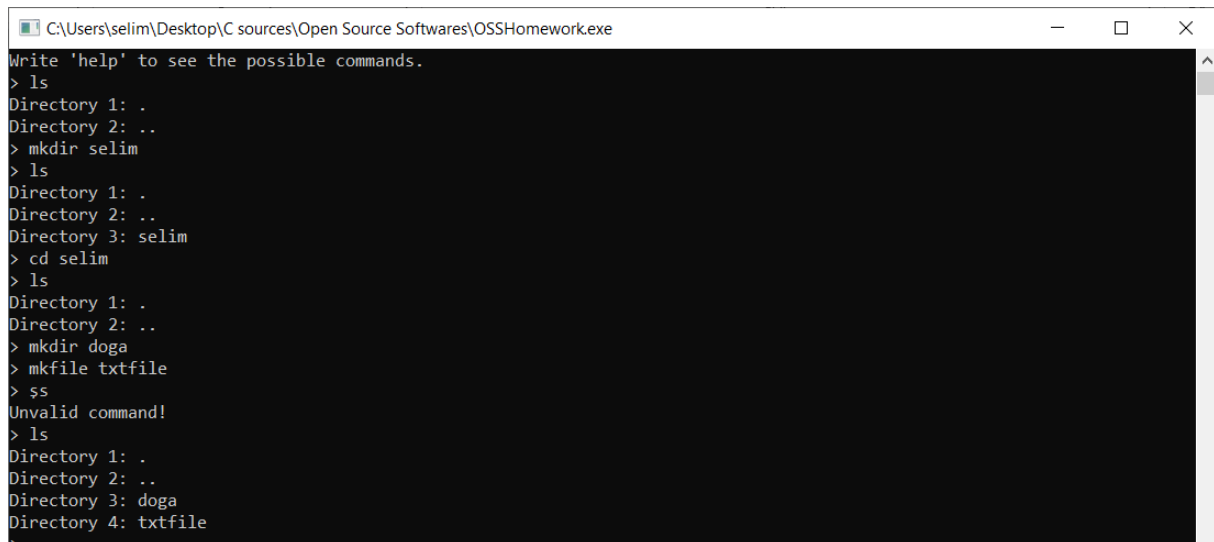
```
C:\Users\selim\Desktop\C sources\Open Source Softwares\OSSHomework.exe
Write 'help' to see the possible commands.
> ls
Directory 1: .
Directory 2: ..
> mkdir selim
> ls
Directory 1: .
Directory 2: ..
Directory 3: selim
>
```

Then, we change the directory to “Selim” with cd command and print what’s inside it with again “ls”.

A screenshot of a terminal window titled "C:\Users\selim\Desktop\C sources\Open Source Softwares\OSSHomework.exe". The terminal shows the prompt "Write 'help' to see the possible commands." followed by the command "> ls". The output is "Directory 1: ." and "Directory 2: ..". Then, the command "> mkdir selim" is entered, followed by "> ls". The output now includes "Directory 3: selim". Then, the command "> cd selim" is entered, followed by "> ls". The output is "Directory 1: ." and "Directory 2: ..". The prompt ">" is shown again.

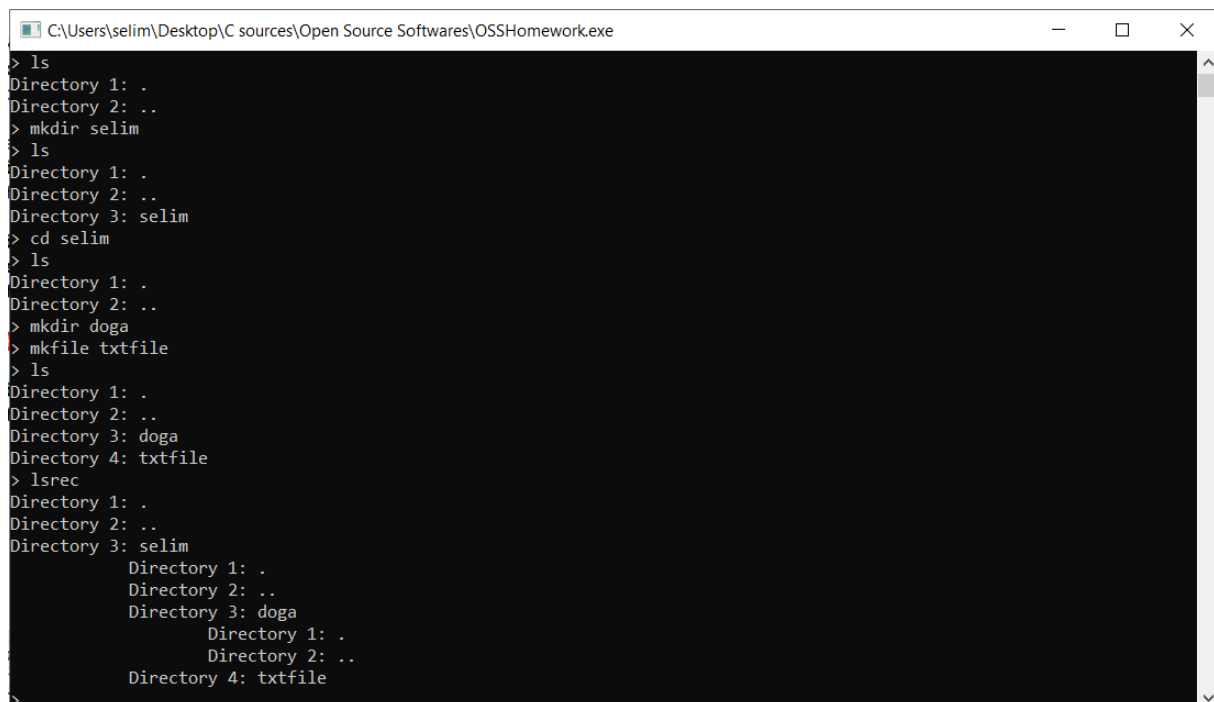
```
C:\Users\selim\Desktop\C sources\Open Source Softwares\OSSHomework.exe
Write 'help' to see the possible commands.
> ls
Directory 1: .
Directory 2: ..
> mkdir selim
> ls
Directory 1: .
Directory 2: ..
Directory 3: selim
> cd selim
> ls
Directory 1: .
Directory 2: ..
>
```

After that, we create another directory and a regular file inside “Selim”. It says “Unvalid command” if we miss spell it (Oops!).



```
C:\Users\selim\Desktop\C sources\Open Source Softwares\OSSHomework.exe
Write 'help' to see the possible commands.
> ls
Directory 1: .
Directory 2: ..
> mkdir selim
> ls
Directory 1: .
Directory 2: ..
Directory 3: selim
> cd selim
> ls
Directory 1: .
Directory 2: ..
> mkdir doga
> mkfile txtfile
> $s
Unvalid command!
> ls
Directory 1: .
Directory 2: ..
Directory 3: doga
Directory 4: txtfile
>
```

Lastly, we write “lsrec” to print the whole file system.



```
C:\Users\selim\Desktop\C sources\Open Source Softwares\OSSHomework.exe
> ls
Directory 1: .
Directory 2: ..
> mkdir selim
> ls
Directory 1: .
Directory 2: ..
Directory 3: selim
> cd selim
> ls
Directory 1: .
Directory 2: ..
> mkdir doga
> mkfile txtfile
> ls
Directory 1: .
Directory 2: ..
Directory 3: doga
Directory 4: txtfile
> lsrec
Directory 1: .
Directory 2: ..
Directory 3: selim
    Directory 1: .
    Directory 2: ..
    Directory 3: doga
        Directory 1: .
        Directory 2: ..
    Directory 4: txtfile
>
```